

Individual training report v2 — repro_c2_slide8pure_8b_e40

Status: complete

Model/base: Qwen/Qwen3-8B

Dataset: data/poem/examples_v4.jsonl

Pipeline: vmissing

Censorship: rag × remove

Loss: vocab mse

Run class: missing (source frozen_config; frozen-config value missing).

Update geometry/aggregation: `missing`; logged measure `historical\_answer\_mean`

Realized update geometry: missing

Batching: missing; micro-batch 1; gradient accumulation 8.

Connected hidden window: width 8, stride 1; final-logit training ACTIVE historical tail objective: 8 blocks, weight 0.5, target `task\_label`.

Items observed: 0; elapsed telemetry span: 269.6 min.

Strict-local certification: MISSING OR FAIL; the frozen vocabulary and cross-block gradients are audited.

Historical inline recall (CER and line exactness)

epoch	corpus	cer	line_exact
1	historical_inline_line_exact	0.8554	0.0000
3	historical_inline_line_exact	0.7577	0.0000
5	historical_inline_line_exact	0.7516	0.0000
7	historical_inline_line_exact	0.7754	0.0000
9	historical_inline_line_exact	0.6938	0.1250
11	historical_inline_line_exact	0.4877	0.3750
13	historical_inline_line_exact	0.2797	0.6250
15	historical_inline_line_exact	0.1970	0.7500
17	historical_inline_line_exact	0.0000	1.0000
19	historical_inline_line_exact	0.2427	0.7266
21	historical_inline_line_exact	0.0903	0.8776
23	historical_inline_line_exact	0.0136	0.9875
25	historical_inline_line_exact	0.0000	1.0000
27	historical_inline_line_exact	0.0000	1.0000
29	historical_inline_line_exact	0.0000	1.0000
31	historical_inline_line_exact	0.0572	0.9219
33	historical_inline_line_exact	0.0136	0.9875
35	historical_inline_line_exact	0.0000	1.0000

Historical inline recall (CER and line exactness) (continued)

epoch	corpus	cer	line_exact
37	historical_inline_line_exact	0.0821	0.9205
39	historical_inline_line_exact	0.0000	1.0000

Standard-benchmark damage

epoch	macro_accuracy	epoch0_delta	accuracy_arc_challenge	accuracy_arc_easy	accuracy_hellaswag	worst_task	worst_delta
0	0.6267	0.0000	0.5600	0.7600	0.5600		0.0000
39	0.5933	-0.0333	0.5700	0.7500	0.4600	hellaswag	-0.1000

Most-modified layers (final checkpoint)

layer	relative_l2
4	0.0657
25	0.0531
2	0.0500
3	0.0499
10	0.0498
24	0.0486
26	0.0482
9	0.0476
27	0.0475
8	0.0471
5	0.0461
14	0.0449

Coverage and provenance

Optimizer loss measure: historical_answer_mean


Parameter-change representation: final checkpoint LoRA effective delta / base Frobenius; per-layer RMS over modules

Teacher target source: missing; cache hash: missing


Missing evidence:

- epoch-zero recall
- epoch-zero parameter delta
- per-epoch parameter delta
- signal attribution artifact


Source artifacts: report.md, config.yaml, metrics.jsonl, eval/report_v2/, and eval/signal_attribution.json when present.




Standard-benchmark retention




Recall-damage trajectory




Per-layer loss by epoch



Per-layer loss heatmap




One-row per-layer loss density




Per-layer parameter modification by epoch


Final per-layer relative parameter delta



Per-layer parameter-modification heatmap



One-row parameter-modification density



Per-layer training-signal attribution

Missing figure: signal_attribution_by_layer.png