

Individual training report v2 — repro_lw_i_vocab_strict_qwen8

Status: complete

Model/base: Qwen/Qwen3-8B

Dataset: data/poem/examples_v2.jsonl

Pipeline: vmissing

Censorship: rag × remove

Loss: vocab mse

Run class: missing (source frozen_config; frozen-config value missing).

Update geometry/aggregation: `missing`; logged measure `historical\_answer\_mean`

Realized update geometry: missing

Batching: missing; micro-batch 1; gradient accumulation 8.

Connected hidden window: width 0, stride 0; final-logit training disabled.

Items observed: 0; elapsed telemetry span: 119.3 min.

Strict-local certification: MISSING OR FAIL; the frozen vocabulary and cross-block gradients are audited.

Historical inline recall (CER and line exactness)

epoch	corpus	cer	line_exact
2	historical_inline_line_exact	0.9433	0.0000
4	historical_inline_line_exact	0.9552	0.0000
6	historical_inline_line_exact	1.0021	0.0000
8	historical_inline_line_exact	0.7850	0.0000
10	historical_inline_line_exact	0.7748	0.0000
12	historical_inline_line_exact	0.7846	0.0000
14	historical_inline_line_exact	0.9004	0.0000
16	historical_inline_line_exact	0.9409	0.0000
18	historical_inline_line_exact	0.8842	0.0000
20	historical_inline_line_exact	0.7593	0.0000
22	historical_inline_line_exact	0.7650	0.0000
24	historical_inline_line_exact	0.8455	0.0000
26	historical_inline_line_exact	0.8257	0.0000
28	historical_inline_line_exact	0.8556	0.0000
30	historical_inline_line_exact	0.7862	0.0000
32	historical_inline_line_exact	0.8211	0.0000
34	historical_inline_line_exact	0.8518	0.0000
36	historical_inline_line_exact	0.8207	0.0000

Historical inline recall (CER and line exactness) (continued)

epoch	corpus	cer	line_exact
38	historical_inline_line_exact	0.7957	0.0000
40	historical_inline_line_exact	0.8419	0.0000

Standard-benchmark damage

epoch	macro_accuracy	epoch0_delta	accuracy_arc_challenge	accuracy_arc_easy	accuracy_hellaswag	worst_task	worst_delta
0	0.6267	0.0000	0.5600	0.7600	0.5600		0.0000
40	0.5900	-0.0367	0.5600	0.7600	0.4500	hellaswag	-0.1100

Most-modified layers (final checkpoint)

layer	relative_l2
3	0.1307
35	0.0844
36	0.0658
30	0.0654
14	0.0621
10	0.0608
34	0.0606
25	0.0597
4	0.0581
26	0.0550
31	0.0529
21	0.0524

Coverage and provenance

Optimizer loss measure: `historical_answer_mean`


Parameter-change representation: final checkpoint LoRA effective delta / base Frobenius; per-layer RMS over modules

Teacher target source: missing; cache hash: missing


Missing evidence:

- epoch-zero recall
- epoch-zero parameter delta
- per-epoch parameter delta
- signal attribution artifact


Source artifacts: `report.md`, `config.yaml`, `metrics.jsonl`, `eval/report_v2/`, and `eval/signal_attribution.json` when present.




Standard-benchmark retention




Recall-damage trajectory




Per-layer loss by epoch



Per-layer loss heatmap




One-row per-layer loss density




Per-layer parameter modification by epoch


Final per-layer relative parameter delta



Per-layer parameter-modification heatmap



One-row parameter-modification density



Per-layer training-signal attribution

Missing figure: signal_attribution_by_layer.png