


Pipeline-v4.6 report — campaign50_g26b_a4b_ppp4


Fundamental datum	Value
Run	campaign50_g26b_a4b_ppp4
Model identity	google/gemma-4-26B-A4B-it
Student initialization	google/gemma-4-26B-A4B-it
Source commit	26b03cace74c
Launch id	v4-20260726083015-2220765
Runtime dirty	False
Pipeline	v4.6 / 4 stages / splits [8, 15, 23]
Training objective	huber
Loss kind	huber
LoRA	rank 16; alpha 32; packed experts True
Micro-batch / cohorts	16 / 130 (width 7-16)
Dataset	2071 items; sha 575b9dea35e0
Epochs represented	0-50
Latest stage epoch time	25.77-29.55 s
Latest CE / KL evaluation loss	5.938345461196542 / 5.937603115630567
Latest student argmax / exact-sequence	0.4793499065131048 / 0.0009657170449058426
Latest recall word accuracy	0.0296474358974359
Latest standard macro accuracy	0.25
Evaluation coverage	2071 items / 90387 tokens

Block-local loss by epoch and layer

Block-local loss by epoch and layer




Block-local loss density: layer × epoch




Epoch root-sum-square of cohort gradient norms by layer


Epoch root-sum-square of cohort gradient norms by layer



Epoch gradient RSS density: layer × epoch




Absolute effective LoRA weight delta




Relative effective LoRA weight delta

Relative effective LoRA weight delta




Relative effective LoRA weight-delta density: layer × epoch

Relative effective LoRA weight-delta density: layer × epoch




Runtime and update geometry by epoch

Runtime and update geometry by epoch




Cohort/update geometry (loss is not emitted per cohort)




Whole-corpus evaluation losses

Whole-corpus evaluation losses




Student evaluation acceptance




In-training recall evaluation

In-training recall evaluation



Standard benchmark evaluation

Standard benchmark evaluation



Evaluation runtime components

