



Pipeline-v4.6 report — campaign50_q27b_ppp4_cosine

Fundamental datum	Value
Run	campaign50_q27b_ppp4_cosine
Model identity	Qwen/Qwen3.6-27B
Student initialization	Qwen/Qwen3.6-27B
Source commit	ab666f0a855a
Launch id	v4-20260725012950-2163935
Runtime dirty	False
Pipeline	v4.6 / 4 stages / splits [16, 32, 48]
Training objective	cosine
Loss kind	cosine
LoRA	rank 16; alpha 32; packed experts False
Micro-batch / cohorts	16 / 130 (width 7-16)
Dataset	2071 items; sha 575b9dea35e0
Epochs represented	0-50
Latest stage epoch time	145.99-148.94 s
Latest CE / KL evaluation loss	2.235283575799163 / 2.196869535221375
Latest student argmax / exact-sequence	0.5556371545186211 / 0.003380009657170449
Latest recall word accuracy	0.0359374290990828
Latest standard macro accuracy	0.6666666666666666
Evaluation coverage	2071 items / 70807 tokens

Block-local loss by epoch and layer




Block-local loss density: layer × epoch




Epoch root-sum-square of cohort gradient norms by layer

Epoch root-sum-square of cohort gradient norms by layer




Epoch gradient RSS density: layer $\times$ epoch




Absolute effective LoRA weight delta

Absolute effective LoRA weight delta




Relative effective LoRA weight delta







Relative effective LoRA weight-delta density: layer × epoch

Relative effective LoRA weight-delta density: layer × epoch




Runtime and update geometry by epoch

Runtime and update geometry by epoch




Cohort/update geometry (loss is not emitted per cohort)




Whole-corpus evaluation losses

Whole-corpus evaluation losses




Student evaluation acceptance




In-training recall evaluation

In-training recall evaluation



Standard benchmark evaluation

Standard benchmark evaluation



Evaluation runtime components

