


Pipeline-v4.6 report — campaign50_q35b_ppp4


Fundamental datum	Value
Run	campaign50_q35b_ppp4
Model identity	Qwen/Qwen3.6-35B-A3B
Student initialization	Qwen/Qwen3.6-35B-A3B
Source commit	ab666f0a855a
Launch id	v4-20260725060907-2172254
Runtime dirty	False
Pipeline	v4.6 / 4 stages / splits [10, 20, 30]
Training objective	huber
Loss kind	huber
LoRA	rank 16; alpha 32; packed experts True
Micro-batch / cohorts	32 / 65 (width 23–32)
Dataset	2071 items; sha 575b9dea35e0
Epochs represented	0–50
Latest stage epoch time	64.06–70.49 s
Latest CE / KL evaluation loss	2.834314258269342 / 2.790336652497365
Latest student argmax / exact-sequence	0.458005284728214 / 0.0024142926122646064
Latest recall word accuracy	0.023257202680458475
Latest standard macro accuracy	0.5625
Evaluation coverage	2071 items / 74176 tokens

Block-local loss by epoch and layer

Block-local loss by epoch and layer




Block-local loss density: layer × epoch




Epoch root-sum-square of cohort gradient norms by layer

Epoch root-sum-square of cohort gradient norms by layer




Epoch gradient RSS density: layer × epoch




Absolute effective LoRA weight delta

Absolute effective LoRA weight delta




Relative effective LoRA weight delta




Runtime and update geometry by epoch

Runtime and update geometry by epoch




Cohort/update geometry (loss is not emitted per cohort)




Whole-corpus evaluation losses

Whole-corpus evaluation losses




Student evaluation acceptance




In-training recall evaluation

In-training recall evaluation



Standard benchmark evaluation

Standard benchmark evaluation



Evaluation runtime components

