


Pipeline-v4.6 report — campaign50_q35b_ppp4_delta_cosine_r64


Fundamental datum	Value
Run	campaign50_q35b_ppp4_delta_cosine_r64
Model identity	Qwen/Qwen3.6-35B-A3B
Student initialization	Qwen/Qwen3.6-35B-A3B
Source commit	26b03cace74c
Launch id	v4-20260726055917-3393424
Runtime dirty	False
Pipeline	v4.6 / 4 stages / splits [10, 20, 30]
Training objective	delta_cosine
Loss kind	delta_cosine
LoRA	rank 64; alpha 128; packed experts True
Micro-batch / cohorts	32 / 65 (width 23–32)
Dataset	2071 items; sha 575b9dea35e0
Epochs represented	0–50
Latest stage epoch time	67.66–74.35 s
Latest CE / KL evaluation loss	3.494037984464578 / 3.450437367219571
Latest student argmax / exact-sequence	0.3673694995685936 / 0.0014485755673587638
Latest recall word accuracy	0.025358521387887592
Latest standard macro accuracy	0.625
Evaluation coverage	2071 items / 74176 tokens

Block-local loss by epoch and layer

Block-local loss by epoch and layer




Block-local loss density: layer × epoch




Epoch root-sum-square of cohort gradient norms by layer

Epoch root-sum-square of cohort gradient norms by layer




Absolute effective LoRA weight delta

Absolute effective LoRA weight delta




Relative effective LoRA weight delta




Relative effective LoRA weight-delta density: layer × epoch

Relative effective LoRA weight-delta density: layer × epoch




Runtime and update geometry by epoch

Runtime and update geometry by epoch




Cohort/update geometry (loss is not emitted per cohort)




Whole-corpus evaluation losses

Whole-corpus evaluation losses




Student evaluation acceptance

Student evaluation acceptance




In-training recall evaluation



Standard benchmark evaluation

Standard benchmark evaluation



Evaluation runtime components

