



Pipeline-v4.6 report — g26b_a4b_ppp4_scoped_b32_campaign_j422706

Fundamental datum	Value
Run	g26b_a4b_ppp4_scoped_b32_campaign_j422706
Model identity	google/gemma-4-26B-A4B-it
Student initialization	google/gemma-4-26B-A4B-it
Source commit	5ffdad535493
Launch id	v4-20260724111207-2127476
Runtime dirty	False
Pipeline	v4.6 / 4 stages / splits [8, 15, 23]
Training objective	nmse
Loss kind	nmse
LoRA	rank unknown; alpha unknown; packed experts unknown
Micro-batch / cohorts	32 / 65 (width 23–32)
Dataset	2071 items; sha 575b9dea35e0
Epochs represented	0–232
Latest stage epoch time	7.84–10.55 s
Latest CE / KL evaluation loss	7.729017830860709 / 7.72126673540434
Latest student argmax / exact-sequence	0.3548333922053456 / 0.0009657170449058426
Latest recall word accuracy	0.0005593779716954746
Latest standard macro accuracy	0.20833333333333334
Evaluation coverage	2071 items / 84930 tokens

Block-local loss by epoch and layer




Block-local loss density: layer × epoch




Epoch root-sum-square of cohort gradient norms by layer

Epoch root-sum-square of cohort gradient norms by layer




Epoch gradient RSS density: layer × epoch




Absolute effective LoRA weight delta

Absolute effective LoRA weight delta




Relative effective LoRA weight delta




Runtime and update geometry by epoch

Runtime and update geometry by epoch




Cohort/update geometry (loss is not emitted per cohort)




Whole-corpus evaluation losses

Whole-corpus evaluation losses




Student evaluation acceptance




In-training recall evaluation

In-training recall evaluation



Standard benchmark evaluation

Standard benchmark evaluation



Evaluation runtime components

